

**Exercise for MA-INF 2213 Computer Vision SS21**  
**25.06.2021**  
**Submission on 06.07.2021**

**1. Bag-of-Words model (5 Points)**

Implement a Bag-of-Words (BoW) model [1] following the subsequent steps:

- (a) Extract SIFT keypoints and features [3] for all images.
- (b) Build a BoW dictionary of 54 words by grouping the features of the previous step (only the *training set*) using K-Means.
- (c) Extract for each image the histogram that corresponds to the BoW dictionary of the previous step.
- (d) *Hint: Read/Write your BoW directory from/to disc to save some time.*

A dataset<sup>1</sup> of 265 images containing 53 object classes is provided in the directory `./images`. Each image contains only one object instance and can be thus labeled with its name. The dataset contains 5 images per class. The first image should be used just for testing (*testing set*), while only the last 4 images should be used for training (*training set*).

Since feature extraction is not the main focus of the exercise, the following (or similar) `OpenCV` structures and methods can be used for this step (Exercise 1): `cv2.BOWKMeansTrainer`, `cv2.FlannBasedMatcher`, `cv2.BOWImgDescriptorExtractor`, `cv2.xfeatures2d.SIFT_create`.

In case the library `cv2.xfeatures2d` is not found, reinstall `python-opencv` via `conda install -c conda-forge opencv` if you use Anaconda. If you use `pip`, install `opencv` via `pip install opencv-python==4.5.1.48`  
`pip install OPENCV-CONTRIB-PYTHON=4.5.1.48`

**2. k-Nearest-Neighbor Classification (5 Points)**

Classify the *testing set* images of the dataset using the SIFT BoW model of the previous step and a *k-Nearest-Neighbor* classifier with a Euclidean distance metric. Compute and print the percentage of false assignments for  $k = 1$  and  $k = 3$ .

**3. Metric Learning (10 Points)**

Implement a classifier for the provided dataset following the approach of [2].

- (a) Based on similarity<sup>2</sup> define all the similar ( $y_{ij} = 1$ ) and dissimilar ( $y_{ij} = 0$ ) pairs of the training set.
- (b) Compute the covariance matrices of the above mentioned sets (Equations 12, 13 of [2]):

$$\Sigma_{y_{ij}=1} = \frac{1}{N_{y_{ij}=1}} \sum_{y_{ij}=1} (\mathbf{x}_i - \mathbf{x}_j)(\mathbf{x}_i - \mathbf{x}_j)^T \quad (1)$$

$$\Sigma_{y_{ij}=0} = \frac{1}{N_{y_{ij}=0}} \sum_{y_{ij}=0} (\mathbf{x}_i - \mathbf{x}_j)(\mathbf{x}_i - \mathbf{x}_j)^T \quad (2)$$

where  $(i, j)$  define a pair of images,  $\mathbf{x}$  is a feature vector and  $N$  denotes the cardinality of the specified set (*similar* and *dissimilar* image pairs).

---

<sup>1</sup>A modified version of the dataset [http://www.vision.ee.ethz.ch/datasets\\_extra/Obj\\_DB.tar.gz](http://www.vision.ee.ethz.ch/datasets_extra/Obj_DB.tar.gz)

<sup>2</sup>A similar pair is a pair of images of the same object from a different viewpoint.

- (c) Learn a squared Mahalanobis distance metric using (non-numbered equation that follows Equation 16 in [2]):

$$\hat{\mathbf{M}} = \Sigma_{y_{ij}=1}^{-1} - \Sigma_{y_{ij}=0}^{-1} \quad (3)$$

- (d) Verify that the learned metric is positive semi-definite by checking if all eigenvalues are positive. In the opposite case recover a positive semi-definite metric  $\mathbf{M}$  by re-projecting  $\hat{\mathbf{M}}$  onto the cone of positive semi-definite matrices by:
- Replacing all negative eigenvalues with a small positive value (0.001).
  - Using the formula

$$M = P \begin{bmatrix} \lambda_1 & 0 & \dots & 0 \\ 0 & \lambda_2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \lambda_n \end{bmatrix} P^{-1} \quad (4)$$

where  $\mathbf{M}$  is the positive semi-definite matrix to be reconstructed,  $P$  is a matrix whose columns are the eigenvectors  $\mathbf{v}_1, \dots, \mathbf{v}_n$  and  $\lambda_1, \dots, \lambda_n$  are the corresponding eigenvalues.

**Hint:** You can use the method `np.linalg.eig` of `numpy` for eigenanalysis.

- (e) Classify the *testing set* images of the dataset using the SIFT BoW model of the first step and a *k-Nearest-Neighbor* classifier with the learned (squared) Mahalanobis distance metric (Equation 16 of [2]):

$$d_M^2(\mathbf{x}_i, \mathbf{x}_j) = (\mathbf{x}_i - \mathbf{x}_j)^T M (\mathbf{x}_i - \mathbf{x}_j) \quad (5)$$

Compute and print the percentage of false assignments for  $k = 1$  and  $k = 3$ .

You are not allowed to use any library apart from the ones provided in the template.

## References

- [1] Gabriella Csurka, Christopher R. Dance, Lixin Fan, Jutta Willamowski, and Cédric Bray. Visual categorization with bags of keypoints. In *In Workshop on Statistical Learning in C.Vision, ECCV*, 2004.
- [2] Martin Koestinger, Martin Hirzer, Paul Wohlhart, Peter M. Roth, and Horst Bischof. Large scale metric learning from equivalence constraints. In *CVPR*, 2012.
- [3] D.G. Lowe. Object recognition from local scale-invariant features. In *ICCV*, 1999.